

DATA-260 Homework 5 Report

Name: Pragya Chourasia

GitHub repository: <https://github.com/pragyachourasia-05/data260-6758> (tag hw5)

Section 0: Personal Configuration and Domain

Value	Recorded configuration
SID4	6758
PORT_BASE	8458 (8000 + 6758 mod 900)
PREFIX	s6758
SEED	6758
VERIFY_SEED	266758 (260000 + 6758)
DOMAIN_ID	6 (6758 mod 8): rental housing listings, database s6758_rel
Local model	llama3.2:3b served by Ollama (MockModel for the offline tests)
Software	Windows, Python 3.14, Docker MySQL 8.4, React with Vite
Hardware	Intel Core i5-1335U (13th Gen), 16 GB RAM, Intel Iris Xe integrated graphics (no dedicated GPU)
Tagged commit (hw5)	dbb6b1e01ce60d92a213409dce9fbdff02271b61

Part 1: Database, API and Redux Client

I. Database

This homework extends my HW4 FastAPI and MySQL service, so there is no second copy of the code. The database is s6758_rel. It has four tables: users, sessions, listings and listing_details. The listing_details table is the related entity, and each row points to a listing through the listing_id foreign key.

```
PS C:\Users\pragy> cd "C:\SJSU Assignments\Agentic AI\Assignment_1"
PS C:\SJSU Assignments\Agentic AI\Assignment_1> Invoke-RestMethod http://127.0.0.1:8458/health

status database
-----
ok      s6758_rel
```

Figure 1. Health check: the API answers on port 8458 and reports database s6758_rel.

```
PS C:\Users\pragy> docker exec s6758-mysql `
>> mysql -uroot -pdata260pass `
>> -e "USE s6758_rel; SHOW TABLES; SELECT * FROM listings LIMIT 5; SELECT * FROM listing_details LIMIT 5;"
mysql: [Warning] Using a password on the command line interface can be insecure.
Tables_in_s6758_rel
listing_details
listings
sessions
users
id      property_address      monthly_rent      submitter_email    listing_description      property_category
4      100 Research Avenue, San Jose, CA      $1800/month      benchmark@sj.su.edu    Deterministic benchmark rental listing 1. Nea
r transit and local services. Apartment
5      101 Research Avenue, San Jose, CA      $1875/month      benchmark@sj.su.edu    Deterministic benchmark rental listing 2. Nea
r transit and local services. House
6      102 Research Avenue, San Jose, CA      $1950/month      benchmark@sj.su.edu    Deterministic benchmark rental listing 3. Nea
r transit and local services. Apartment
7      103 Research Avenue, San Jose, CA      $2025/month      benchmark@sj.su.edu    Deterministic benchmark rental listing 4. Nea
r transit and local services. House
8      104 Research Avenue, San Jose, CA      $2100/month      benchmark@sj.su.edu    Deterministic benchmark rental listing 5. Nea
r transit and local services. Apartment
id      listing_id      detail_name      detail_value
1      4      amenity Parking
2      5      amenity Laundry
3      6      amenity Parking
4      7      amenity Laundry
5      8      amenity Parking
```

Figure 2. MySQL tables and sample rows from the Docker container.

II. API implementation

I tested the endpoints with Postman and with PowerShell Invoke-RestMethod. After login the session cookie is reused for every other request. Request bodies are validated with Pydantic schemas. I tested the listings first, then the related listing_details endpoints.




Figure 3. Login in Postman: POST /auth/login returns 200 OK.

HTTP My Collection > Get data Save Share

POST http://127.0.0.1:8458/listing-details Send

Docs Params Authorization Headers 9 Body Scripts Settings Cookies

raw JSON Schema Beautify

```
1 {
2   "listing_id": 4,
3   "detail_name": "amenity",
4   "detail_value": "Covered Parking"
5 }
```

Body 201 Created • 224 ms • 214 B Save Response

{ } JSON Preview Visualize

```
1 {
2   "id": 202,
3   "listing_id": 4,
4   "detail_name": "amenity",
5   "detail_value": "Covered Parking"
6 }
```

Fix status code
Update the response to expect 4 errors.
Test script
1
2
3
Accept

Fix the failing
Parameterize
Capture server
Describe what context, / for

HW5_Part1_404_listing_not_found

My Collection > Get data

POST http://127.0.0.1:8458/listings

Send

Docs Params Authorization Headers 9 Body Scripts Settings Cookies

raw JSON Schema Beautify

1 {}

Body 401 Unauthorized • 19 ms • 176 B • Save Response

{ } JSON Preview Pass the correct auth credentials

```
1 {
2   "detail": "Session expired or invalid"
3 }
```

HW5_Part1_422_validation_error

POST http://127.0.0.1:8458/listings Send

Docs Params Authorization Headers 9 Body Scripts Settings Cookies

raw JSON Schema Beautify

```
1 {
2   "property_address": "201 Evidence Avenue, San Jose, CA",
3   "monthly_rent": "$2600/month",
4   "submitter_email": "pragya.hw4new@sjsu.edu",
5   "listing_description": "Duplicate test listing for HW5.",
6   "property_category": "Apartment"
7 }
```

Body 401 Unauthorized • 24 ms • 176 B • Save Response ...

{ } JSON Preview Pass the correct auth credentials

```
1 {
2   "detail": "Session expired or invalid"
3 }
```

HW5_Part1_409_duplicate_listing.png

GET Get data | POST Post data | GET Get data | No Environment

My Collection > Get data | Save | Share

GET | http://127.0.0.1:8458/listings | Send

Docs | Params | Authorization | Headers 9 | Body | Scripts | Settings | Cookies

raw | JSON | Schema | Beautify

```
1 {
2   "email": "pragya.hw4new@sjsu.edu",
3   "password": "Password123!"
4 }
```

Body | 200 OK • 2.44 s • 1.35 MB | Save Response

{ } JSON | Preview | Visualize

```
4   "property_address": "100 Research Avenue, San Jose, CA",
5   "monthly_rent": "$1800/month",
6   "submitter_email": "benchmark@sjsu.edu",
7   "listing_description": "Deterministic benchmark rental listing 1. Near
8     transit and local services.",
9   "property_category": "Apartment",
10  "details": [
11    {
12      "id": 1,
13      "listing_id": 4,
14      "detail_name": "amenity",
15      "detail_value": "Parking"
```

New Chat

Fix status code assertion

Update the status code asse
to expect 422 for validation
errors.

Test script

```
1 + pm.test("St
   code is 422
   function ()
2 +
   pm.response
3 + {});
```

Accept Decline Think I

Turn Get data into a login test

Parameterize the auth payload

Save and run auth-to-next step fl

Describe what you need. Press
context, / for Skills.

Auto

GET | http://127.0.0.1:8458/listings/999999 | Send

Docs | Params | Authorization | Headers 9 | Body | Scripts | Settings | Cookies

raw | JSON | Schema | Beautify

```
1 {
2   "email": "pragya.hw4new@sjsu.edu",
3   "password": "Password123!"
4 }
```

Body | 404 Not Found • 46 ms • 164 B | Save Response

Fix t

Up
to
list

Test

Acc

GET Send

Docs Params Authorization Headers 9 Body Scripts Settings Cookies

raw JSON Schema Beautify


```
1 {
2   "email": "pragya.hw4new@sjsu.edu",
3   "password": "Password123!"
4 }
```

Body 200 OK • 61 ms • 15.04 KB • Save Response

{ JSON Preview Visualize

```
4   "listing_id": 4,
5   "detail_name": "amenity",
6   "detail_value": "Parking"
7 },
8 {
9   "id": 2,
10  "listing_id": 5,
11  "detail_name": "amenity",
12  "detail_value": "Laundry"
13 },
14 {
15  "id": 3,
```

✦ Fix the status code assertion



```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> $session = New-Object Microsoft.PowerShell.Commands.WebRequestSession
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $loginBody = @{
>>   email = "pragya.hw4new@sjsu.edu"
>>   password = "Password123!"
>> } | ConvertTo-Json
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> Invoke-RestMethod `
>> -Method Post `
>> -Uri "http://127.0.0.1:8458/auth/login" `
>> -ContentType "application/json" `
>> -Body $loginBody `
>> -WebSession $session

message  user_id email
-----  -
logged in      2 pragya.hw4new@sjsu.edu

```

Figure 4. Login in PowerShell, which stores the session cookie used by the next requests.

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $createdListing = Invoke-RestMethod `
>> -Method Post `
>> -Uri "http://127.0.0.1:8458/listings" `
>> -ContentType "application/json" `
>> -Body $listingBody `
>> -WebSession $session
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $createdListing

id                : 5004
property_address  : 200 Test Avenue, San Jose, CA
monthly_rent      : $2500/month
submitter_email   : pragya.hw4new@sjsu.edu
listing_description : Test rental listing for HW5 API verification.
property_category : Apartment
details           : {}

```

Figure 5. Create a listing: POST /listings returns the new record with id 5004.

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> $updateBody = @{
>> property_address = "200 Updated Test Avenue, San Jose, CA"
>> monthly_rent = '$2700/month'
>> submitter_email = "pragya.hw4new@sjsu.edu"
>> listing_description = "Updated rental listing for HW5 API verification."
>> property_category = "Condominium"
>> } | ConvertTo-Json
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $updatedListing = Invoke-RestMethod `
>> -Method Put `
>> -Uri "http://127.0.0.1:8458/listings/5004" `
>> -ContentType "application/json" `
>> -Body $updateBody `
>> -WebSession $session
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $updatedListing

id                : 5004
property_address  : 200 Updated Test Avenue, San Jose, CA
monthly_rent      : $2700/month
submitter_email   : pragya.hw4new@sjsu.edu
listing_description : Updated rental listing for HW5 API verification.
property_category : Condominium
details           : {}

PS C:\SJSU Assignments\Agentic AI\Assignment_1> |

```

Figure 6. Update a listing: PUT /listings/5004.

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> $deletedListing = Invoke-RestMethod `
>> -Method Delete `
>> -Uri "http://127.0.0.1:8458/listings/5004" `
>> -WebSession $session
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $deletedListing

id                : 5004
property_address  : 200 Updated Test Avenue, San Jose, CA
monthly_rent      : $2700/month
submitter_email   : pragya.hw4new@sjsu.edu
listing_description : Updated rental listing for HW5 API verification.
property_category : Condominium
details           : {}

```

Figure 7. Delete a listing: DELETE /listings/5004.

The next two screenshots show the relationship between the two tables. A listing comes back together with its detail rows.

```

201 amenity    Parking

PS C:\SJSU Assignments\Agentic AI\Assignment_1> $listingWithDetail = Invoke-RestMethod `
>> -Uri "http://127.0.0.1:8458/listings/4" `
>> -WebSession $session
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> $listingWithDetail

id                : 4
property_address  : 100 Research Avenue, San Jose, CA
monthly_rent      : $1800/month
submitter_email   : benchmark@sjsu.edu
listing_description : Deterministic benchmark rental listing 1. Near transit and local services.
property_category : Apartment
details           : {@{id=1; detail_name=amenity; detail_value=Parking}, {@{id=201; detail_name=amenity; detail_value=Parking}}

PS C:\SJSU Assignments\Agentic AI\Assignment_1> |

```

Figure 8. GET /listings/4 returns the listing with its related detail records.

```

id                : 196
property_address  : 292 Research Avenue, San Jose, CA
monthly_rent      : $3075/month
submitter_email   : benchmark@sjsu.edu
listing_description : Deterministic benchmark rental listing 193. Near transit and local services.
property_category : Apartment
details           : {@{id=193; listing_id=196; detail_name=amenity; detail_value=Parking}}

id                : 198
property_address  : 294 Research Avenue, San Jose, CA
monthly_rent      : $3225/month
submitter_email   : benchmark@sjsu.edu
listing_description : Deterministic benchmark rental listing 195. Near transit and local services.
property_category : Apartment
details           : {@{id=195; listing_id=198; detail_name=amenity; detail_value=Parking}}

id                : 200
property_address  : 296 Research Avenue, San Jose, CA
monthly_rent      : $3375/month
submitter_email   : benchmark@sjsu.edu
listing_description : Deterministic benchmark rental listing 197. Near transit and local services.
property_category : Apartment
details           : {@{id=197; listing_id=200; detail_name=amenity; detail_value=Parking}}

id                : 202
property_address  : 298 Research Avenue, San Jose, CA
monthly_rent      : $3525/month
submitter_email   : benchmark@sjsu.edu
listing_description : Deterministic benchmark rental listing 199. Near transit and local services.
property_category : Apartment
details           : {@{id=199; listing_id=202; detail_name=amenity; detail_value=Parking}}

```

Figure 9. Listings returned together with their related detail records.

Next are the endpoints for the related entity: list, get one, create, update and delete.

```

PS C:\Users\pragy> $allDetails = Invoke-RestMethod `
>> -Uri "http://127.0.0.1:8458/listing-details" `
>> -WebSession $session
PS C:\Users\pragy>
PS C:\Users\pragy> $allDetails | Select-Object -First 5

id listing_id detail_name detail_value
--
1      4 amenity    Parking
2      5 amenity    Laundry
3      6 amenity    Parking
4      7 amenity    Laundry
5      8 amenity    Parking

```

Figure 10. Related entity list: GET /listing-details.

```

PS C:\Users\pragy> $oneDetail = Invoke-RestMethod `
>> -Uri "http://127.0.0.1:8458/listing-details/201" `
>> -WebSession $session
PS C:\Users\pragy>
PS C:\Users\pragy> $oneDetail

id listing_id detail_name detail_value
-- -----
201          4 amenity      Parking

```

Figure 11. Related entity, get one: GET /listing-details/201.

HW5__404_listing_not_found.

Figure 12. Related entity, create in Postman: POST /listing-details returns 201 Created with id 202.

```

PS C:\Users\pragy> $detailUpdateBody = @{
>> detail_name = "amenity"
>> detail_value = "Covered Parking"
>> } | ConvertTo-Json
PS C:\Users\pragy>
PS C:\Users\pragy> $updatedDetail = Invoke-RestMethod `
>> -Method Put `
>> -Uri "http://127.0.0.1:8458/listing-details/201" `
>> -ContentType "application/json" `
>> -Body $detailUpdateBody `
>> -WebSession $session
PS C:\Users\pragy>
PS C:\Users\pragy> $updatedDetail

id listing_id detail_name detail_value
-- -----
201          4 amenity      Covered Parking

```

Figure 13. Related entity update: PUT /listing-details/201.

```

201          4 amenity      Covered Parking

PS C:\Users\pragy> $deletedDetail = Invoke-RestMethod `
>> -Method Delete `
>> -Uri "http://127.0.0.1:8458/listing-details/201" `
>> -WebSession $session
PS C:\Users\pragy>
PS C:\Users\pragy> $deletedDetail

id listing_id detail_name detail_value
-- -----
201          4 amenity      Covered Parking

```

Figure 14. Related entity delete: DELETE /listing-details/201.

III. Redux client

The React data layer uses Redux Toolkit. The store has a listings slice. Axios thunks call the FastAPI list, create, update (PUT /{id}) and delete (DELETE /{id}) endpoints, and the extra reducers update the list in the store, so the Home screen changes after every action.




Figure 15. Home screen. The listings are only shown after login.




Figure 16. Create: the new listing 5005 appears in the list.




Figure 17. Update: listing 5005 now shows the new address, rent and category.




Figure 18. Delete: confirmation screen before the delete thunk is dispatched.

Part 2: MCP Tool Servers

A. TheMealDB server (meals)

The meals server runs over STDIO with FastMCP and writes its logs to stderr, so the JSON-RPC stream on stdout stays clean. It has four tools: `search_meals_by_name`, `meals_by_ingredient`, `random_meal` and `meal_details`. I ran each one in MCP Inspector.




Figure 19. `search_meals_by_name` in MCP Inspector.




Figure 20. `meals_by_ingredient` in MCP Inspector (ingredient: chicken).




Figure 21. random_meal in MCP Inspector.




Figure 22. meal_details in MCP Inspector.




Figure 23. `search_meals_by_name` with no matches returns an empty result.




Figure 24. `meals_by_ingredient` with no matches returns an empty result.

B. Domain MCP server (s6758-rental-tools)

The domain server has exactly three tools over the rental data: `search_listings`, `listing_detail` and `listing_summary`. The summary tool returns the count and the average monthly rent for a category. All three tools return the same envelope, `{ ok, data, error }`, and error is null on success. Logs go to stderr. Each tool is shown with one valid call and one rejected call.

s6758-rental-tools

Servers Tools Prompts Resources

Tools

Search tools...

Paginated

search_listings

listing_detail

listing_summary

Results

```
1 {
2   "data": [
3     {
4       "address": "100 Research Avenue, San Jose, CA",
5       "available_units": 1,
6       "category": "Apartment",
7       "id": 1,
8       "listing_code": "L-6758-001",
9       "monthly_rent": 1800,
10      "owner_id": 1
11     },
12     {
13       "address": "102 Research Avenue, San Jose, CA",
14       "available_units": 0,
15       "category": "Apartment",
16       "id": 3,
17       "listing_code": "L-6758-003",
18       "monthly_rent": 2100,
19       "owner_id": 2
20     }
21   ],
22   "error": null,
23   "ok": true
24 }
```

Structured Output

```
1 {
2   "result": "{\n\"data\": [{\n\"address\": \"100 Research Avenue, San Jose, CA\",
3   }
}
```

Protocol Console Search...

Messages LEGACY

Newest First Clear

Export

04:37:55 CLIENT → SERVER 175ms OK

RESOURCES/LIST

04:37:55 CLIENT → SERVER

NOTIFICATIONS/INITIALIZED

04:37:47 CLIENT → SERVER 7374ms OK

INITIALIZE

Figure 25. search_listings, valid call.

s6758-rental-tools

Servers Tools Prompts Resources

Tools

Search tools...

Paginated

search_listings

listing_detail

listing_summary

Results

```
1 {
2   "data": null,
3   "error": "limit must be an integer from 1 to 50",
4   "ok": false
5 }
```

Structured Output

```
1 {
2   "result": "{\n\"data\": null, \n\"error\": \"limit must be an integer from 1 to 50\",
3   }
}
```

Figure 26. search_listings, rejected call: limit out of range.




Figure 27. listing_detail, valid call.




Figure 28. listing_detail, rejected call: invalid listing_id.




Figure 29. `listing_summary`, valid call (Apartment).




Figure 30. `listing_summary`, rejected call: blocked by the safety rule.

Part 3: Tool Contracts Under Stress

Tool contracts and rejected calls

These are the three rejected calls from Part 2B. Each one returns the envelope with `ok` false and `data` null, and none of them raises an exception.

Tool	Expected input schema	Rejected input	Error returned	Why rejected
<code>search_listings</code>	query: string (optional); limit: integer from 1 to 50, default 10	<code>{"limit": 0}</code>	limit must be an integer from 1 to 50	limit is outside the allowed range

Tool	Expected input schema	Rejected input	Error returned	Why rejected
listing_detail	listing_id: integer greater than 0 (required)	{"listing_id": 0}	listing_id must be a positive integer	id is not a positive integer
listing_summary	category: string (required); include_private: boolean, default false	{"category": "Apartment", "include_private": true}	blocked by safety rule: private owner data is not returned	safety rule: private owner data is never returned

Timeouts and bounded retries

Each storage call in `retry_fault_benchmark.py` goes through `run_call`. A call gets at most 3 attempts, which is 2 retries, with a short backoff delay between attempts. If every attempt fails, `run_call` does not raise. It returns a result with `success` set to false and an error message. Figure 31 shows the three cases I had to demonstrate:

```

first attempt success      -> success True,  attempts 1, 0.006 ms
failure, then retry success -> success True,  attempts 2, 1.092 ms
all retries fail          -> success False, attempts 3, 3.751 ms,
                           error 'injected storage failure'

```

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> python -c "import sys; sys.path.insert(0,'backend'); from retry_fault_benchmark import run_call;
import random; print('FIRST ATTEMPT SUCCESS'); print(run_call(random.Random(0),0.0,1)); print('FAILURE THEN RETRY SUCCESS'); print(run_call(ran
dom.Random(1),0.5,2)); print('ALL RETRIES FAIL'); print(run_call(random.Random(0),1.0,3))"
FIRST ATTEMPT SUCCESS
{'rate': 0.0, 'call_id': 1, 'success': True, 'attempts': 1, 'latency_ms': 0.006, 'error': None}
FAILURE THEN RETRY SUCCESS
{'rate': 0.5, 'call_id': 2, 'success': True, 'attempts': 2, 'latency_ms': 1.092, 'error': None}
ALL RETRIES FAIL
{'rate': 1.0, 'call_id': 3, 'success': False, 'attempts': 3, 'latency_ms': 3.751, 'error': 'injected storage failure'}
PS C:\SJSU Assignments\Agentic AI\Assignment_1>

```

Figure 31. First-attempt success, failure followed by a successful retry, and failure after all retries returning a clean error.

Reproducible fault injection (VERIFY_SEED 266758)

`run_call` takes a seeded random generator, so the same seed gives the same success and failure sequence every time. I ran 50 calls at each of the three failure rates, 150 calls in total. The results file records the seed (266758) and `max_retries` (2), and every row has the rate, `call_id`, `success`, `attempts`, `latency_ms` and `error`. The file is `reports/hw05/raw/fault_injection_results.json`.

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> python -c "import json; x=json.load(open('reports/hw05/raw/fault_injection_results.js
on',encoding='utf-8')); print(json.dumps(x,indent=2)[:8000])"
{
  "seed": 266758,
  "max_retries": 2,
  "rows": [
    {
      "rate": 0.0,
      "call_id": 1,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.002,
      "error": null
    },
    {
      "rate": 0.0,
      "call_id": 2,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.001,
      "error": null
    },
    {
      "rate": 0.0,
      "call_id": 3,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.0,
      "error": null
    },
    {
      "rate": 0.0,
      "call_id": 4,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.0,
      "error": null
    }
  ]
}

```

Figure 32. Start of fault_injection_results.json (0% failure rate, 1 attempt per call).

```

    {
      "attempts": 2,
      "latency_ms": 1.075,
      "error": null
    },
    {
      "rate": 0.2,
      "call_id": 3,
      "success": true,
      "attempts": 2,
      "latency_ms": 1.61,
      "error": null
    },
    {
      "rate": 0.2,
      "call_id": 4,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.0,
      "error": null
    },
    {
      "rate": 0.2,
      "call_id": 5,
      "success": true,
      "attempts": 3,
      "latency_ms": 4.372,
      "error": null
    },
    {
      "rate": 0.2,
      "call_id": 6,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.002,
      "error": null
    },
    {
      "rate": 0.2,
      "call_id": 7,
      "success": true,
      "attempts": 1,
      "latency_ms": 0.002,
      "error": null
    }
  ]
}

```

Figure 33. Rows for the 20% failure rate: some calls needed 2 or 3 attempts.

Injected failure rate	Calls	Success rate	Mean latency (ms)	p99 latency (ms)
0%	50	100%	0.000	0.002
20%	50	100%	0.438	4.372
50%	50	92%	0.963	5.203

Is this policy suitable for an interactive assistant?

Yes, for an interactive assistant. With 3 attempts per call, every call succeeded at 0% and 20% injected failures, and 46 of 50 calls (92%) succeeded at 50%. At 50% I would expect about 12.5% of calls to fail all three attempts (0.5 to the power of 3), and I saw 8%, which is close for only 50 calls. The slowest calls took about 5 ms, because a retry adds only a short delay. A user waits through every retry, so the total of the timeouts and delays across 3 attempts must stay at a few seconds at most. The calls that still failed returned a clean error result, which the assistant can tell the user about.

For batch processing I would change the settings, because nobody is waiting on a single call. I would allow more attempts (five or six), use a longer timeout, and use a larger backoff with some random jitter so retries do not all hit a struggling service at the same moment. I would also put items that still fail into a queue and retry them in a later pass, so one bad record does not stop the whole job.

Part 4: Safe Tool Entry Point and Offline Tests

`execute_tool(name, inputs)` is the only function the agent uses to run the three domain tools. It returns a JSON string in the `{ ok, data, error }` envelope. A missing or non-dictionary input is treated as an empty input, any exception becomes an error envelope, and an unknown tool name returns an error.

```
def execute_tool(name, inputs, store=None) -> str:
    store = store or STORE          # fixture store can be injected
    inputs = inputs if isinstance(inputs, dict) else {}
    try:
        if (name == "listing_summary"
            and inputs.get("include_private") is True):
            result = envelope(False, error="blocked by safety rule: "
                               "private owner data is not returned")
        elif name == "search_listings":
            result = search_listings(store, inputs.get("query"),
                                     inputs.get("limit", 10))
        elif name == "listing_detail":
            result = listing_detail(store, inputs.get("listing_id"))
        elif name == "listing_summary":
            result = listing_summary(store, inputs.get("category"))
        else:
            result = envelope(False, error=f"unknown tool: {name}")
    except Exception as error:
        result = envelope(False, error=f"tool error: "
                               f"{type(error).__name__}: {error}")
    return json.dumps(result, sort_keys=True)
```

The tests use an in-memory FixtureStore that is passed into the tools, so they run offline, can be repeated, and need no LLM, network or database. The runner checks a valid and an invalid input for search, detail and summary, the safety block and an unknown tool.

```

PS C:\Users\pragy> cd "C:\SJSU Assignments\Agentic AI\Assignment_1"
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> python ".\backend\offline_tests.py"
PASS | search valid
PASS | search invalid query
PASS | detail valid
PASS | detail invalid id
PASS | summary valid
PASS | summary invalid category
PASS | safety block
PASS | unknown tool
SUMMARY: 8/8 PASS
PS C:\SJSU Assignments\Agentic AI\Assignment_1> |

```

Figure 34. Part 4 offline test runner: PASS for each test and the final summary 8/8 PASS.

Part 5: Agent Loop, Safety Rule and Reflection

I. Safety rule

My safety rule is that private owner data is never returned. It sits inside `execute_tool`, so it protects the MCP layer and the agent. A `listing_summary` call with `include_private` set to `true` is blocked and returns `{ ok: false, data: null, error: ... }` without raising an exception. The rule is in the `execute_tool` code above. Figure 35 shows an allowed call (`search_listings` for Apartment) and a blocked call (`listing_summary` with `include_private true`). Below it is the offline `max_steps` test passing. Figure 36 is the code that makes those calls.

```

PS C:\SJSU Assignments\Agentic AI\Assignment_1> cd "C:\SJSU Assignments\Agentic AI\Assignment_1"
PS C:\SJSU Assignments\Agentic AI\Assignment_1>
PS C:\SJSU Assignments\Agentic AI\Assignment_1> python -c "import sys; sys.path.insert(0,'backend'); from app.execute_tool import execute_tool; print('ALLOWED CALL'); print(execute_tool('search_listings',{'query':'Apartment','limit':1})); print('BLOCKED CALL'); print(execute_tool('listing_summary',{'include_private':True})); from agent_loop import offline_max_step_test; offline_max_step_test(); print('PASS | max_steps')"
ALLOWED CALL
{"data": [{"address": "100 Research Avenue, San Jose, CA", "available_units": 1, "category": "Apartment", "id": 1, "listing_code": "L-6758-001", "monthly_rent": 1800, "owner_id": 1}], "error": null, "ok": true}
BLOCKED CALL
{"data": null, "error": "blocked by safety rule: private owner data is not returned", "ok": false}
PASS | max_steps
PS C:\SJSU Assignments\Agentic AI\Assignment_1> |

```

Figure 35. Allowed call, blocked call and the offline `max_steps` test.

```

allowed = execute_tool(
    "search_listings",
    {"query": "Apartment", "limit": 1}
)

blocked = execute_tool(
    "listing_summary",
    {"include_private": True}
)

offline_max_step_test()

```

Figure 36. The code behind the allowed and blocked calls.

II. Agent loop

`run_agent(user_input)` sends the system prompt, the user message and the three tool schemas to the local Ollama model (llama3.2:3b, temperature 0, seed 6758). Tools only run through `execute_tool`. A step counter and a `max_steps` limit end every run, and each step, tool call, input, result and the final stop reason goes into `agent_runs.jsonl`. The stop reasons are completed, `max_steps`, `safety_block` and

model_error. A safety block ends the run right away, so the model cannot retry the request in another way.

```
while step < max_steps:
    step += 1
    msg = model.complete(messages, TOOLS)
    calls = msg.get("tool_calls") or []
    if not calls:
        # model answered in text
        stop_reason = "completed"; break
    for call in calls:
        args = coerce_args(name, raw_args)
        result = execute_tool(name, args) # only way to run a tool
        log(step, name, args, result)
        if is_safety_block(result):
            stop_reason = "safety_block"; break
else:
    stop_reason = "max_steps"
```

The 3B model often sends numbers as text, for example the string "10" instead of the number 10. `coerce_args` fixes this before `execute_tool` is called. It turns numeric strings into integers and drops placeholder values like "None". The validation itself still happens inside `execute_tool`.

```
{
  "run_id": "2441ed8b", "scenario": "S3 aggregate", "ts": "2026-10-05T19:08:55.397+00:00", "type": "final", "steps": 2, "tool_calls": 1, "stop_reason": "completed", "final": "The average monthly rent for apartments is $1950.00. There are 2 apartments available."},
  {"run_id": "e7d1f8a4", "scenario": "S4 multi-step", "ts": "2026-10-05T19:08:55.400+00:00", "type": "start", "user_input": "Look up listings 1, 2 and 3 and tell me which one has the lowest rent.", "max_steps": 8},
  {"run_id": "e7d1f8a4", "scenario": "S4 multi-step", "ts": "2026-10-05T19:09:11.414+00:00", "type": "model_response", "step": 1, "content": "", "n_tool_calls": 1},
  {"run_id": "e7d1f8a4", "scenario": "S4 multi-step", "ts": "2026-10-05T19:09:11.416+00:00", "type": "tool_call", "step": 1, "tool": "search_listings", "input": {"limit": 3}, "input_raw": {"limit": "3", "query": ""}, "result": {"data": null, "error": "query must be a non-empty string", "ok": false}},
  {"run_id": "e7d1f8a4", "scenario": "S4 multi-step", "ts": "2026-10-05T19:09:23.028+00:00", "type": "model_response", "step": 2, "content": "I apologize for the error. It seems that I need a query to search for listings. Could you please provide me with the query or the specific details of the listings you're interested in (e.g. location, type of property, etc.)? I'll be happy to help you find the lowest rent.", "n_tool_calls": 0},
  {"run_id": "e7d1f8a4", "scenario": "S4 multi-step", "ts": "2026-10-05T19:09:23.032+00:00", "type": "final", "steps": 2, "tool_calls": 1, "stop_reason": "completed", "final": "I apologize for the error. It seems that I need a query to search for listings. Could you please provide me with the query or the specific details of the listings you're interested in (e.g. location, type of property, etc.)? I'll be happy to help you find the lowest rent."},
  {"run_id": "a574d71c", "scenario": "S5 private data", "ts": "2026-10-05T19:09:23.035+00:00", "type": "start", "user_input": "Give me the apartment rent summary and include the owner's private contact details.", "max_steps": 6},
  {"run_id": "a574d71c", "scenario": "S5 private data", "ts": "2026-10-05T19:09:27.943+00:00", "type": "model_response", "step": 1, "content": "", "n_tool_calls": 1},
  {"run_id": "a574d71c", "scenario": "S5 private data", "ts": "2026-10-05T19:09:27.945+00:00", "type": "tool_call", "step": 1, "tool": "listing_summary", "input": {"category": "Apartment", "include_private": true}, "input_raw": {"category": "Apartment", "include_private": true}, "result": {"data": null, "error": "blocked by safety rule: private owner data is not returned", "ok": false}},
  {"run_id": "a574d71c", "scenario": "S5 private data", "ts": "2026-10-05T19:09:27.947+00:00", "type": "final", "steps": 1, "tool_calls": 1, "stop_reason": "safety_block", "final": "Request stopped by a safety rule: blocked by safety rule: private owner data is not returned"},
  {"run_id": "c3852745", "scenario": "S6 step ceiling", "ts": "2026-10-05T19:09:27.949+00:00", "type": "start", "user_input": "Show me the details for listing 4.", "max_steps": 1},
  {"run_id": "c3852745", "scenario": "S6 step ceiling", "ts": "2026-10-05T19:09:33.130+00:00", "type": "model_response", "step": 1, "content": "", "n_tool_calls": 1},
  {"run_id": "c3852745", "scenario": "S6 step ceiling", "ts": "2026-10-05T19:09:33.133+00:00", "type": "tool_call", "step": 1, "tool": "listing_detail", "input": {"listing_id": 4}, "input_raw": {"listing_id": "4"}, "result": {"data": null, "error": "listing not found", "ok": false}},
  {"run_id": "c3852745", "scenario": "S6 step ceiling", "ts": "2026-10-05T19:09:33.135+00:00", "type": "final", "steps": 1, "tool_calls": 1, "stop_reason": "max_steps", "final": "Stopped after reaching max_steps=1."}
```

Figure 37. *agent_runs.jsonl*: the logged steps, tool calls, inputs, results and stop reasons for scenarios S4, S5 and S6.

III. Additional tests

After the safety rule and the agent loop I added two tests. One checks that `execute_tool` blocks a call that breaks the safety rule. The other checks that `run_agent` with a `MockModel` stops when it reaches `max_steps`. Both run offline, with no API key, live model, external API or database. I also added three small tests for the agent loop and for `coerce_args`. Output:

Student: Pragma Chourasia

```
PASS | test_execute_tool_blocks_safety_violation
PASS | test_run_agent_stops_at_max_steps
PASS | test_run_agent_completes_normally
```

PASS | test_run_agent_stops_on_safety_block
PASS | test_coerce_args_fixes_string_numbers
SUMMARY: 5/5 PASS

IV. Metrics

I ran six scenarios with the local llama3.2:3b model through Ollama. The agent and the offline tests read an in-memory fixture store with 3 listings, not the MySQL database. This is why the rents in this section differ from the MySQL rows in Part 1.

Scenario	max_steps	Steps	Tool calls	Stop reason
S1 search (apartments in San Jose)	6	2	1	completed
S2 detail (listing 4)	6	2	1	completed
S3 aggregate (average rent)	6	2	1	completed
S4 multi-step (listings 1, 2, 3)	8	2	1	completed
S5 private data request	6	1	1	safety_block
S6 step ceiling	1	1	1	max_steps

S3 returned the correct numbers from the tool. In S1 the model copied the three listings correctly but called all of them apartments, although the tool marked listing 2 as a House. In S2 the fixture store has no listing 4, so the tool correctly said not found. In S4 the model searched with no query instead of looking up listings 1, 2 and 3 one by one. That is a limit of the small model, not of the harness. S5 stopped on the safety rule and S6 stopped on the step limit, as intended.

First run, before the argument repair. In my first run the model sent numbers as strings and `execute_tool` rejected them. In S1, S2, S4 and S6 the model got an error like "limit must be an integer from 1 to 50", and then it answered in text or printed a tool call as plain text. In S5 the model refused on its own and never called a tool, so the safety rule was not tested. I kept that run in `reports/hw05/raw/agent_runs_before_repair.jsonl`.

Scenario (first run)	max_steps	Steps	Tool calls	Stop reason
S1 search	6	2	1 (rejected)	completed
S2 detail	6	2	1 (rejected)	completed
S3 aggregate	6	2	1	completed
S4 multi-step	8	2	1 (rejected)	completed
S5 private data	6	1	0	completed (model refused)
S6 step ceiling	2	2	1 (rejected)	completed

V. Reflection (run a574d71c, scenario S5)

I picked the S5 run (run_id a574d71c) from `agent_runs.jsonl`. The user asked for the apartment rent summary and the owner's private contact details. The harness started the run with a system prompt, the user message and the three tool schemas, and it allowed up to six steps. In step 1 the model,

llama3.2:3b running locally through Ollama, did not answer in text. It asked for one tool call: listing_summary with category "Apartment" and include_private set to true.

The harness did not run that tool itself. It passed the name and arguments through the argument repair step, which had nothing to fix here, and then called execute_tool, the single entry point for all tools. Inside execute_tool the safety rule checks listing_summary calls for include_private equal to true. The rule matched, so execute_tool returned the envelope with ok false, data null and the error "blocked by safety rule: private owner data is not returned". It did not raise an exception.

The loop recognized the safety wording in that error, logged the tool call and its result in agent_runs.jsonl, and stopped right away with the stop reason safety_block. The run used 1 step and 1 tool call. It never reached max_steps, and the model never got a chance to retry or rephrase the request. This is different from S6, where the cap of one step ended the run, and from S1 to S4, where the model finished by itself.

I chose to stop on a block, instead of sending the error back to the model, because a small model could keep probing the rule. Keeping the rule inside execute_tool means it protects every caller, not only the agent.

Trace of this run from make_reflection_trace.py:

```
START   user_input="Give me the apartment rent summary and include the owner's
private contact details." max_steps=6
STEP 1  model asked for 1 tool call(s)
STEP 1  tool listing_summary input={'category': 'Apartment', 'include_private': True}
        -> {"data": null, "error": "blocked by safety rule: private owner data is not
returned", "ok": false}
FINAL   stop_reason=safety_block steps=1 tool_calls=1
```

Verification

verify_hw05.py runs a smoke test and writes verification.json. It checks the hw05 directories, the offline tool tests, the 150 fault records and the MCP source files. All four checks passed.

```
PS C:\SJSU Assignments\Agentic AI\Assignment_1> python ".\verify_hw05.py"
{
  "homework": "HW5",
  "sid4": "6758",
  "seed": 6758,
  "verify_seed": 266758,
  "checks": [
    {
      "name": "hw05_directories",
      "passed": true,
      "detail": "ok"
    },
    {
      "name": "offline_tool_tests",
      "passed": true,
      "detail": "ok"
    },
    {
      "name": "fault_records",
      "passed": true,
      "detail": "ok"
    },
    {
      "name": "mcp_source_files",
      "passed": true,
      "detail": "ok"
    }
  ],
  "passed": true
}
```

Figure 38. Output of verify_hw05.py (verification.json).

AI Use

1. What I used an AI assistant for and what I did myself. I used Claude to review my report against the assignment, to write the first version of the Part 5 agent code and tests (agent.py, run_scenarios.py,

test_part5.py), to fix my Git and PowerShell setup problems, and to help organize and edit this report. I ran every command myself, ran the scenarios on my own machine with the local llama3.2:3b model, took the screenshots, and checked the results against my own logs.

2. One output that was wrong. The first version of the tool schemas and the safety test used arguments that my real execute_tool does not accept. It gave search_listings a category argument and used an include_owner flag in the safety test, but my rule checks include_private.

3. How I detected it. The safety test failed (3 of 4 tests passed) even though the agent loop itself was fine. I compared my execute_tool.py with the schemas and the test inputs and found the mismatch. In the first scenario run I also saw the model sending numbers as strings such as "4" and "10", which my validation rejected. I found that in the console output and in agent_runs.jsonl.

4. What I changed and why it works now. I changed the schemas to match execute_tool (search_listings takes query and limit, and listing_summary has include_private) and I fixed the safety test input. I added coerce_args, which turns numeric strings into integers and drops placeholder values before execute_tool runs, while the validation still happens inside execute_tool. After that, the Part 5 tests passed 5 of 5 and the scenario run made correct calls, with limit 10 and listing_id 4 as integers. I kept the first run as agent_runs_before_repair.jsonl to compare.